

# ■ Reckonsys Gen AI Dev Assignment: Agentic RAG Challenge

**Objective:** Build a basic agentic RAG system where the agent autonomously decides whether to retrieve answers from documents or search the web when confidence is low.

## ■■ Functional Requirements

- 1. Document Ingestion:** Load and index 5–10 text or PDF documents using embeddings + vector store (FAISS, Chroma, or Milvus). Include metadata like title and source.
- 2. Retriever-based QA:** Implement RAG: query → retrieve → generate. Include confidence scoring based on similarity or LLM evaluation.
- 3. Agentic Decision Layer:** Agent decides whether to return RAG output or trigger web search if confidence < threshold.
- 4. Web Search Agent:** Fetch and summarize information from web using DuckDuckGo, Tavily, or SerpAPI. Include note: “■ This answer was found on the web.”
- 5. Output Format:** Return JSON: {"source": "document" | "web", "answer": "string", "confidence": "float", "context\_snippets": ["string"]}

## ■ Tech Expectations

- Backend: Python (FastAPI or Streamlit demo preferred)
- LLM: OpenAI GPT-4o-mini / GPT-4-Turbo or open-source
- Libraries: LangChain, LlamaIndex, FastMCP, Haystack
- Vector DB: FAISS / Chroma / Pinecone
- Search: Tavily / DuckDuckGo / SerpAPI

## ■ Bonus Points

- Implement ReAct-style reasoning or tool-calling flow.
- Add UI or CLI interface.
- Include logs or reasoning trace.
- Deploy via FastAPI or Gradio.
- Provide README.md with architecture overview.

## ■ Submission Deliverables

1. GitHub repo with code, docs, and sample docs folder.
2. Optional demo video (≤ 2 mins).
3. Deadline: 72 hours from assignment start.

## ■ Evaluation Criteria

Criteria	Weight
Architecture clarity	25%
RAG setup & performance	25%
Agentic logic implementation	25%
Code quality & documentation	15%
Creativity / Bonus features	10%